



Data Management for Data Science  
Sapienza Università di Roma  
2018/2019

# SQL

Maurizio Lenzerini, Riccardo Rosati

Dipartimento di Ingegneria informatica, automatica  
e gestionale Antonio Ruberti  
Sapienza Università di Roma



# SQL

## 1. Data definition

1. **Data definition**
2. Data manipulation
3. Queries
4. Further aspects



# SQL

- SQL ("S**tr**uctured **Q**uery **L**anguage") contains both the DDL (Data Definition Language) and the DML (Data Manipulation Language)
- Different versions of the language exist
- Brief history:
  - First proposal: **SEQUEL** (IBM Research, 1974);
  - first implementations in SQL/DS (IBM) and Oracle (1981);
  - from 1983 ca., “standard de facto”
  - standard versions released in 1986, 1989, **1992**, 1999, 2003
  - Last versions only partially supported by commercial systems



# Using an SQL-based DBMS

- An SQL-based database management system (DBMS) is a **server** that allows for managing a **set** of relational databases
- Following the relational model, an SQL database is characterized by a **schema** (intensional level) and by an **instance** (extensional level)
- In addition, an SQL database is characterized by a set of **meta-data** (catalog)



## Data definition in SQL

- The most important statement of the SQL DDL is  
**create table**
  - Defines a relation schema (specifying attributes and constraints)
  - Creates an empty instance of the relation schema
- Syntax:  

```
create table TableName (  
    AttributeName Domain [ Constraints ]  
    .....  
    AttributeName Domain [ Constraints ]  
    [ OtherConstraints]  
)
```

## create table: example

```
create table Employee (  
  ID          character(6) primary key,  
  Name        character(20) not null,  
  Surname     character(20) not null,  
  Depart      character(15),  
  Salary      numeric(9) default 0,  
  City        character(15),  
  foreign key (Depart) references  
    Department (DepName),  
  unique (Surname, Name)  
)
```

table  
name

constraint

attribute  
name

domain  
(type)



## SQL and the relational model

- **Remark:** an SQL table is defined as a **multiset** of n-tuples
- Only if the table has a primary key (or a set of attributes defined as unique), the same n-tuple cannot appear twice in the table



# Domains for attributes

- **Predefined domains**

- **Character:**

- `char(n)` or `character(n)`
    - `varchar(n)` (or `char varying(n)`)
    - `nchar(n)` and `nvarchar(n)` (or `nchar varying(n)`) (UNICODE)

- **Numeric:**

- `int` or `integer`, `smallint`
    - `numeric`, (or `numeric(p)`, `numeric(p,s)`)
    - `decimal`, (or `decimal(p)`, `decimal(p,s)`)
    - `float`, `float(p)`, `real`, `double precision`

- **Date, time:**

- `Date`, `time`, `timestamp`
    - `time with timezone`, `timestamp with timezone`

- **Bit:**

- `bit(n)`
    - `bit varying(n)`

- **Further domains (introduced in SQL:1999)**

- `boolean`
    - `BLOB`, `CLOB`, `NCLOB` (binary/character large object)





# Domains for attributes

- **User-defined domains**

- Sintassi

```
create domain NewDomainName  
as PreExistingDomain [ Default ] [ Constraints ]
```

- *Example:*

```
create domain Grade  
as smallint default null  
check ( value >=18 and value <= 30 )
```



## Intra-relational constraints

- **not null** (over single attributes)
- **unique**: defines a set of attributes as a super-key:
  - single attribute:  
**unique** after the domain specification
  - Multiple attributes:  
**unique** (*Attribute*, ..., *Attribute*)
- **primary key**: (only one primary key can be defined on a relation) syntax similar to **unique**; implies **not null**
- **check**, for more complex constraints



## Example

```
create table Employee (  
    ID            character(6) primary key,  
    Name          character(20) not null,  
    Surname       character(20) not null,  
    Depart        character(15) ,  
    Salary        numeric(9) default 0,  
    City          character(15) ,  
    foreign key (Depart) references  
        Department (DepName) ,  
    unique (Surname, Name)  
)
```



## primary key, alternative

```
create table Employee (  
    ID character(6) primary key,  
    ...  
)
```

oppure

```
create table Employee (  
    ID character(6) ,  
    ...  
    primary key (ID)  
)
```



## Keys over multiple attributes

```
create table Employee ( ...  
    Name      character(20) not null,  
    Surname   character(20) not null,  
    unique (surname,name)  
)
```

is **different** from:

```
create table Employee ( ...  
    Name      character(20) not null unique,  
    Surname   character(20) not null unique  
)
```



# Inter-relational constraints

- **check**, for complex constraints
- **references** and **foreign key** allow for defining **referential** integrity constraints

## Syntax:

- single attribute:

**references** after the specification of the domain

- multiple attributes:

**foreign key (Attribute, ..., Attribute) references ...**

The attributes referenced in the end table must constitute a key (**primary key** or **unique**). If they are missing, the attributes of the primary key are considered.

**Semantics**: every combination (without NULL) of values for the attributes in the starting table must appear in the end table



# Inter-relational constraints: example

## Infringements

<u>Code</u>	Date	Policeman	Prov	Number
34321	1/2/95	3987	MI	39548K
53524	4/3/95	3295	TO	E39548
64521	5/4/96	3295	PR	839548
73321	5/2/98	9345	PR	839548

## Policemen

<u>ID</u>	Surname	Name
3987	Rossi	Luca
3295	Neri	Piero
9345	Neri	Mario
7543	Mori	Gino

# Inter-relational constraints: example (cont.)

## Infringements

<u>Code</u>	Date	Policeman	Prov	Number
34321	1/2/95	3987	MI	39548K
53524	4/3/95	3295	TO	E39548
64521	5/4/96	3295	PR	839548
73321	5/2/98	9345	PR	839548

## Car

<u>Prov</u>	<u>Number</u>	Surname	Name
MI	39548K	Rossi	Mario
TO	E39548	Rossi	Mario
PR	839548	Neri	Luca





## Inter-relational constraints: example

```
create table Infringements (  
    Code          character(6) not null primary key,  
    Date          date not null,  
    Policeman     integer not null  
                    references Policemen(ID) ,  
    Prov          character(2) ,  
    Number        character(6) ,  
    foreign key(Prov, Number)  
                    references Car(Prov,Number)  
)
```



# Schema modification : alter table

**alter table**: allows for modifying a table

*Example:*

```
create table Infringements (  
    Code      character(6) not null primary key,  
    Date       date not null,  
    Policeman  integer not null  
                references Policemen(ID) ,  
    Prov       character(2) ,  
    Number     character(6) ,  
)
```

```
alter table Infringements  
add constraint MyConstraint foreign key(Prov, Number)  
references Car(Prov, Number)
```

It can be used to realize **cyclic** referential integrity constraints



# Schema modification: drop table

**drop table**: eliminates a table

Syntax:

```
drop table TableName restrict | cascade
```

*Esempio:*

```
drop table Infringements restrict or simply
```

```
drop table Infringements
```

– eliminates the table if it is not referenced

```
drop table Infringements cascade
```

– eliminates the table and all the tables (and the other database objects) referring to it



## Definition of indices

- Is very important for the system performance
- Deals with the physical level of the DB, not the logical one
- **create index**
- Syntax (simplified):  
**create [unique] index *IndexName* on  
*TableName Attribute,...,Attribute***
- *Example:*  
**create index IndiceIP on  
Infringements (Prov)**



## Catalog (or data dictionary)

Every DBMS creates and maintains special tables that collect the meta-data about

- **tables**
- **attributes**
- ...

For instance, the **Columns** table contains the attributes

- **Column\_Name**
- **Table\_name**
- **Ordinal\_Position**
- **Column\_Default**
- ...



# SQL

## 2. Data manipulation

1. Data definition
- 2. Data manipulation**
3. Queries
4. Further aspects



# Update operations in SQL

- Update operations:
  - addition: `insert`
  - elimination: `delete`
  - modification: `update`
- Of one or multiple tuples of a relation
- Based on a condition that may involve the relation and/or other relations



## Insert: syntax

```
insert into Table [ ( Attributes ) ]  
      values ( Values )
```

or

```
insert into Table [ ( Attributes ) ]  
      select ...
```





## Insert: example

```
insert into person values ('Mario',25,52)
```

```
insert into person(name, age, income)  
values ('Pino',25,52)
```

```
insert into person(name, income)  
values ('Lino',55)
```

```
insert into person (name)  
select father  
from isFather  
where father not in (select name from person)
```



## Insert: comments

- The order of the attributes and the values (if present) is significant
- The list of attributes and the list of values must have the same number of elements
- If the attribute list is missing, all the attributes of the relation are considered, according to the order in which they have been defined
- If the attribute list does not contain all the attributes of the relation, a null value is inserted for every missing attribute (or a default value, if declared)



# Tuple elimination

Syntax:

```
delete from Table [ where Condition ]
```

*Example:*

```
delete from person  
where age < 35
```

```
delete from isFather  
where child not in  
      (select name from person)
```



## Delete: comments

- Deletes the tuples satisfying the condition
- It may cause (if the referential integrity constraints are defined using **cascade**) deletions in other relations
- remember: if the **where** clause is omitted, it is considered as **where true**



# Tuple modification

- **Syntax:**

**update** *TableName*

**set** *Attribute* = < *Expression* | **select** ... | **null** | **default** >  
[ **where** *Condition* ]

- **Semantics:** the tuples that satisfy the «where» condition are deleted

- *Examples:*

```
update person set income = 45  
where name = 'Piero'
```

```
update person set income = income * 1.1  
where age < 30
```



# SQL

## 3. Queries

1. Data definition
2. Data manipulation
- 3. Queries**
4. Further aspects



# The select statement (basic version)

- The query statement in SQL is  
**select**
- It defines a query and returns the result as a table

```
select   Attribute ... Attribute  
from    Table ... Table  
[where   Condition]
```

- The three sections of the statement are usually called:
  - **target list**
  - **from clause**
  - **where clause**



## isMother

<b>mother</b>	<b>child</b>
Luisa	Maria
Luisa	Luigi
Anna	Olga
Anna	Filippo
Maria	Andrea
Maria	Aldo

## isFather

<b>father</b>	<b>child</b>
Sergio	Franco
Luigi	Olga
Luigi	Filippo
Franco	Andrea
Franco	Aldo

## person

<b>name</b>	<b>age</b>	<b>income</b>
Andrea	27	21
Aldo	25	15
Maria	55	42
Anna	50	35
Filippo	26	30
Luigi	50	40
Franco	60	20
Olga	30	41
Sergio	85	35
Luisa	75	87





# Selection and projection

Name and income of people who are less than 30 years old:

$\text{PROJ}_{\text{name, income}}(\text{SEL}_{\text{age} < 30}(\text{person}))$

```
select person.name, person.income
from   person
where  person.age < 30
```

name	income
Andrea	21
Aldo	15
Filippo	30

## Name conventions

- To avoid ambiguity, every attribute name is composed of

*TableName.AttributeName*

- When there is no ambiguity, *TableName* can be omitted

```
select person.name, person.income
from    person
where   person.age < 30
```

can be written as follows:

```
select name, income
from    person
where   age < 30
```



## SELECT, abbreviations

```
select person.name, person.income
from    person
where   person.age < 30
```

can be also written as:

```
select p.name as name, p.income as income
from    person as p
where   p.age < 30
```

or:

```
select p.name as name, p.income as income
from    person p
where   p.age < 30
```



# Projection

surname and city of all employees

**employees**

ID	surname	city	salary
7309	Neri	Napoli	55
5998	Neri	Milano	64
9553	Rossi	Roma	44
5698	Rossi	Roma	64

**PROJ<sub>surname, city</sub> (employees)**



## Projection and duplicates

```
select surname,  
       city  
from employees
```

surname	city
Neri	Napoli
Neri	Milano
Rossi	Roma
Rossi	Roma

```
select distinct surname,  
       city  
from employees
```

surname	city
Neri	Napoli
Neri	Milano
Rossi	Roma



## SELECT, usage of “as”

“**as**” is used in the attribute list to specify a name for an attribute of the result. If such a name is not specified, then the attribute name of the result is equal to the corresponding attribute of the input table.

*Example:*

```
select name as personName, income as salary
from    person
where   age < 30
```

returns a relation with two attributes: **personName** and **salary**

```
select name, income
from    person
where   age < 30
```

returns a relation with two attributes: **name** and **income**



## Exercise 1

Compute the table obtained from table **person** selecting only the people whose income is between 20 and 30, and adding an attribute that has the same value as the attribute **income** in every tuple

Show the result of the query over the table **person** shown at page 32.

**person**

name	age	income
------	-----	--------

## Solution, Exercise 1

```
select name, age, income,  
       income as repeatedIncome  
from   person  
where  income >= 20 and income <= 30
```

name	age	income	repeatedIncome
Andrea	27	21	21
Filippo	26	30	30
Franco	60	20	20





## Selection, without projection

name, age and income of people who are less than 30 years old

**SEL<sub>age<30</sub>(person)**

```
select *  
from   person  
where  age < 30
```

Is an abbreviation for:

```
select name, age, income  
from   person  
where  age < 30
```



all attributes



## Projection, without selection

name and income of all people:

**PROJ**<sub>name, income</sub>(**person**)

```
select name, income  
from    person
```

Is an abbreviation for:

```
select p.name, p.income  
from    person p  
where   true
```



## Expressions in the target list

```
select income/2 as semesterIncome
from   person
where  name = 'Luigi'
```

## Complex condition in the “where” clause:

```
select *
from   person
where  income > 25
       and (age < 30 or age > 60)
```



## “LIKE” condition

People having a name whose first letter is 'A', and whose third letter is 'd':

```
select *  
from   person  
where  name like 'A_d%'
```



## Null values

Employees whose age is or might be greater than 40:

**SEL** `age > 40 OR age IS NULL` (**employees**)

```
select *  
from   employees  
where  age > 40 or age is null
```



## Exercise 2

Compute the tble obtained from table **employees** selecting only the ones whose city is Roma or Milano, projecting the data on the attribute **salary**, and adding an attribute having, in every tuple, a value that is the double the value of the attribute **salary**

Show the result of the query over the table shown at page 36.

**employees**

ID	surname	city	salary
----	---------	------	--------

## Solution, Exercise 2

```
select salary,  
       salary*2 as doubleSalary  
from   employees  
where  city = 'Milano' or  
       city = 'Roma'
```

salary	doubleSalary
64	128
44	88
64	128



## Selection, projection and join

- the **select** statements with a single relation in the **from** clause allow for expressing:
  - selections
  - projections
  - renamings
- **joins** (and cartesian products) are expressed using two or more relations in the **from** clause





# SQL and relational algebra

Given the relations  $R1(A1,A2)$  and  $R2(A3,A4)$ :

The semantics of the query

```
select R1.A1, R2.A4
from   R1, R2
where  R1.A2 = R2.A3
```

can be described as a combination of:

- cartesian product (**from**)
- selection (**where**)
- projection (**select**)

Remark: this does not mean that the DBMS necessarily computes the cartesian product to answer the query!



## SQL: DBMS execution of queries

- SQL expressions are declarative, and we are describing their semantics
- In practice, DBMSs execute operations in efficient ways, for instance:
  - They execute selections as soon as possible
  - If possible, they execute joins instead of cartesian product
- The ability of DBMSs to optimize queries makes it usually not necessary to deal with efficiency when a query is specified



## SQL and relational algebra, 2

Given the relations  $R1(A1,A2)$  and  $R2(A3,A4)$

```
select R1.A1, R2.A4
from   R1, R2
where  R1.A2 = R2.A3
```

corresponds to :

$$\text{PROJ}_{A1,A4} (\text{SEL}_{A2=A3} (R1 \text{ JOIN } R2))$$



## SQL and relational algebra, 3

Renamings may be necessary:

- in the target list
- In the from clause (cartesian product), in particular when the same table must be referred multiple times

```
select X.A1 as B1, ...  
from    R1 X, R2 Y, R1 Z  
where   X.A2 = Y.A3 and ...
```

can be written as

```
select X.A1 as B1, ...  
from    R1 as X, R2 as Y, R1 as Z  
where   X.A2 = Y.A3 and ...
```



## SQL and relational algebra: example

```
select X.A1 as B1, Y.A4 as B2
from    R1 X, R2 Y, R1 Z
where   X.A2 = Y.A3 and Y.A4 = Z.A1
```

```
RENB1,B2←A1,A4 (
  PROJA1,A4 (SELA2 = A3 and A4 = C1 (
    R1 JOIN R2 JOIN RENC1,C2 ← A1,A2 (R1))))
```



## isMother

<b>mother</b>	<b>child</b>
Luisa	Maria
Luisa	Luigi
Anna	Olga
Anna	Filippo
Maria	Andrea
Maria	Aldo

## Is Father

<b>father</b>	<b>child</b>
Sergio	Franco
Luigi	Olga
Luigi	Filippo
Franco	Andrea
Franco	Aldo

## person

<b>name</b>	<b>age</b>	<b>income</b>
Andrea	27	21
Aldo	25	15
Maria	55	42
Anna	50	35
Filippo	26	30
Luigi	50	40
Franco	60	20
Olga	30	41
Sergio	85	35
Luisa	75	87



## Exercise 3: selection, projection and join

Return the fathers of people who earn more than 20 millions.



## Exercise 3: solution

Return the fathers of people who earn more than 20 millions.

$\text{PROJ}_{\text{father}}(\text{isFather JOIN}_{\text{child=name}} \text{SEL}_{\text{income}>20}(\text{person}))$

```
select distinct isFather.father
from    person, isFather
where   isFather.child = person.name
        and person.income > 20
```





## Exercise 4: join

Return the father and the mother of every person.



## Exercise 4: solution

Return the father and the mother of every person.

This can be expressed in relational algebra through the **natural join**.

isFather JOIN isMother

In SQL:

```
select isMother.child, father, mother
from   isMother, isFather
where  isFather.child = isMother.child
```

## Exercise 4: solution

If we interpret the question as: return father and mother of every person appearing in the «person» table, then we need an additional join:

In relational algebra:

$$\text{PROJ}_{\text{child, father, mother}} ((\text{isMother JOIN isFather}) \\ \text{JOIN}_{\text{child=name}} \text{person})$$

In SQL:

```
select isMother.child, father, mother
from   isMother, isFather, person
where  isFather.child = isMother.child
       and isMother.child = person.name
```



## Exercise 5: join and other operations

Return the persons earning more than their fathers, showing name, income and father's income.

## Exercise 5: solution

Return the persons earning more than their fathers, showing name, income and father's income.

PROJ<sub>name, income, RP</sub> (SEL<sub>income>RP</sub>  
(REN<sub>NP,EP,RP</sub> ← name,age,income (person)  
JOIN<sub>NP=father</sub>  
(isFather JOIN<sub>child =name</sub> person)))

```
select      c.name, c.income, p.income
from        person p, isFather t, person c
where       p.name = t.father and
            t.child = c.name and
            c.income > p.income
```



## SELECT, with renaming of the result

Return the persons earning more than their fathers, showing name, income and father's income.

```
select child, c.income as income,  
       p.income as fatherIncome  
from   person p, isFather t, person c  
where  p.name = t.father and  
       t.child = c.name and  
       c.income > p.income
```



## SELECT with explicit join

```
select ...  
from Table { join Table on JoinCondition }, ...  
[ where OtherCondition ]
```

this is the SQL operator corresponding to theta-join



## Explicit join

Return the father and the mother of every person:

```
select isFather.child, father, mother
from   isMother, isFather
where  isFather.child = isMother.child
```

```
select mother, isFather.child, father
from   isMother join isFather on
       isFather.child = isMother.child
```

explicit  
join





## Exercise 6: explicit join

Return the persons earning more than their fathers, showing name, income and father's income.

Express the query using the explicit join.



## SELECT with explicit join: example

Return the persons earning more than their fathers, showing name, income and father's income.

```
select c.name, c.income, p.income
from   person p, isFather t, person c
where  p.name = t.father and
       t.child = c.name and
       c.income > p.income
```

Using the explicit join:

```
select c.name, c.income, p.income
from   person p join isFather t on p.name=t.father
       join person c on t.child=c.name
where  c.income > p.income
```



## Natural join (less frequently used)

$\text{PROJ}_{\text{child, father, mother}}(\text{isFather JOIN}_{\text{child} \leftarrow \text{name}} \text{REN}_{\text{name} \leftarrow \text{child}}(\text{isMother}))$

In algebra:       $\text{isFather JOIN isMother}$

In SQL (with  
Explicit join):      

```
select isFather.child, father, mother
from   isMother join isFather on
       isFather.child = isMother.child
```

In SQL (with  
natural join) :      

```
select isFather.child, father, mother
from   isMother natural join isFather
```



## Left outer join

Return every pair child/father and, if known, the mother.

```
select isFather.child, father, mother
from   isFather left outer join isMother
       on isFather.child = isMother.child
```

(if the mother does not exist, a null value is returned)

Remark: “outer” is optional

```
select isFather.child, father, mother
from   isFather left join isMother
       on isFather.child = isMother.child
```



## Right outer join

if we use the **right** outer join:

```
select isFather.child, father, mother
from   isFather right outer join isMother
       on isFather.child = isMother.child
```

the query returns **all** mothers (even those who do not have a join with isFather)



## Left and right outer join: examples

```
select isFather.child, father, mother
from   isMother join isFather
       on isMother.child = isFather.child
```

```
select isFather.child, father, mother
from   isMother left outer join isFather
       on isMother.child = isFather.child
```

```
select isFather.child, father, mother
from   isMother right outer join isFather
       on isMother.child = isFather.child
```



## Full outer join: examples

```
select isFather.child, father, mother
from   isMother full outer join isFather
       on isMother.child = isFather.child
```

```
select name, father, mother
from   person full outer join isMother on
       person.name = isMother.child
       full outer join isFather on
       person.name = isFather.child
```



## Ordering the result: order by

name and income of persons whose age is less than 30  
in alphabetical order

```
select name, income  
from person  
where age < 30  
order by name
```



ascending  
order

```
select name, income  
from person  
where age < 30  
order by name desc
```



descending  
order





## Ordering the result: order by

```
select name, income  
from person  
where age < 30
```

name	income
Andrea	21
Aldo	15
Filippo	30

```
select name, income  
from person  
where age < 30  
order by name
```

name	income
Aldo	15
Andrea	21
Filippo	30



## Bounding the size of the query result

```
select name, income
from   person
where  age < 30
order by name
limit 2
```



# Bounding the size of the query result

```
select name, income  
from   person  
where  age < 30  
order by name  
limit 2
```

name	income
Andrea	21
Aldo	15



# Aggregate operators

The target list may contain expressions that compute values based on sets of tuples:

- count, min, max, average, total

(simplified) syntax:

*Function ( [ distinct ] ExpressionOverAttributes )*



# Aggregate operators: count

## Syntax:

- Count the number of tuples:

`count (*)`

- Count the values in an attribute:

`count (Attributo)`

- Count the **distinct** values in an attribute:

`count (distinct Attributo)`



## count: example and semantics

*Esempio*: Return the number of children of Franco:

```
select count(*) as NumChildrenFranco
from   isFather
where  father = 'Franco'
```

**Semantics**: the aggregate operator (**count**) is applied to the result of the following query:

```
select *
from   isFather
where  father = 'Franco'
```

## count: example

isFather

father	child
Sergio	Franco
Luigi	Olga
Luigi	Filippo
Franco	Andrea
Franco	Aldo

**NumChildrenFranco**

2



## count and null values

```
select count(*)  
from person
```

Result = 4

```
select count(income)  
from person
```

Result = 3

```
select count(distinct income)  
from person
```

Result = 2

**person**

name	age	income
Andrea	27	21
Aldo	25	NULL
Maria	55	21
Anna	50	35





## Other aggregate operators

### **sum, avg, max, min**

- Allow an attribute or an expression as argument (not “\*”)
- **sum** and **avg**: numeric or date/time arguments
- **max** and **min**: arguments on which a total ordering is defined

*Esempio*: return the income average of Franco's children:

```
select avg(income)
from    person join isFather on
        name = child
where   father = 'Franco'
```




# Aggregate operators and null values

```
select avg(income) as averageIncome  
from person
```

person

name	age	income
Andrea	27	30
Aldo	25	NULL
Maria	55	36
Anna	50	36



averageIncome
34



## Aggregate operators and target list

The following query does not make sense:

```
select name, max(income)
from person
```

For the query to make sense, the **target list** must be **homogeneous**, for instance:

```
select min(age), avg(income)
from person
```



## Aggregate operators and grouping

- In the previous cases, the aggregate operators were applied to all the tuples constituting the query result
- In many cases, we want the aggregate functions to be applied to **partitions of tuples**
- To specify such partitions, the clause **group by** can be used:

**group by** *AttributeList*



# Aggregate operators and grouping

Return the number of children of every father:

```
select father, count(*) as NumChildren
from   isFather
group by father
```

**isFather**

father	child
Sergio	Franco
Luigi	Olga
Luigi	Filippo
Franco	Andrea
Franco	Aldo



father	NumChildren
Sergio	1
Luigi	2
Franco	2



# Semantics of queries with aggregate operators and grouping

1. execute the query **ignoring the group by clause** and the aggregate operators:

```
select *  
from isFather
```

2. group the **tuples having the same values for the attributes mentioned in the group by clause**
3. the aggregate operator is applied to every group and a tuple is produced for every group



## Exercise 7: group by

Return the maximum income (and the age) of every group of persons whose age is greater than 18 and have the same age.

**person**

name	age	income
------	-----	--------



## Exercise 7: solution

Return the maximum income (and the age) of every group of persons whose age is greater than 18 and have the same age.

```
select age, max(income)
from    person
where   age > 18
group by age
```





## Grouping and target list

In query that uses the **group by** clause, the target list should be «homogeneous», namely, **only** attributes appearing in the **group by** clause and aggregation functions should appear in the list.

### *Example:*

- Income of persons, grouped by age (**non-homogeneous** target list):

```
select age, income
from person
group by age
```

- Average income of persons, grouped by age (**homogeneous**, target list (in every group there is only one average income):

```
select age, avg(income)
from person
group by age
```



## Non-homogeneous target list

What happens if the target list is non-homogeneous?

Some systems do not raise any error and for each group return one of the values associated with the group.

*Example:*

Income of persons, grouped by age:

```
select age, income
from person
group by age
```

The DBMS MySQL, for instance, does not raise any error: for each group, it chooses one of the incomes appearing in the group and returns such a value as the income attribute of the target list.



## Conditions on groups

We can also impose **selection conditions on groups**. Group selection is **obviously different** from the condition that selects the tuples forming the groups (**where** clause).

Group selection is realized by the **having** clause, which must appear after the “**group by**” clause.

*Example:* return the fathers whose children have an average income greater than 25.

```
select father, avg(c.income)
from   person c join isFather
      on child = name
group by father
having avg(c.income) > 25
```



## Exercise 8: where or having?

Return the fathers whose children under 30 have an average income greater than 20.



## Exercise 8: solution

Return the fathers whose children under 30 have an average income greater than 20.

```
select father, avg(c.income)
from    person c join isFather
        on child = name
where   c.age < 30
group by father
having  avg(c.income) > 20
```



## Syntax of select statement (summary)

*SelectSQL ::=*

**select**      *AttributeOrExpressionList*  
**from**        *TableList*  
[ **where**      *SimpleConditions* ]  
[ **group by** *GroupingAttributeList* ]  
[ **having**     *AggregationConditions* ]  
[ **order by** *OrderingAttributeList* ]  
[ **limit**      *Number* ]



## Union, intersection and difference

A single **select** statement does not allow for expressing unions (for instance, the union of two tables)

An explicit statement is needed:

```
select ...  
union [all]  
select ...
```

With **union**, duplicate tuples are eliminated

With **union all**, duplicate tuples are kept



## Positional notation

```
select father, child
from   isFather
union
select mother, child
from   isMother
```

Which are the attributes of the result? It depends on the system:

- new names established by the system
- the names of the first select statement
- ...





## Union: result

father	child
Sergio	Franco
Luigi	Olga
Luigi	Filippo
Franco	Andrea
Franco	Aldo
Luisa	Maria
Luisa	Luigi
Anna	Olga
Anna	Filippo
Maria	Andrea
Maria	Aldo



## Positional notation: example

```
select father, child
from isFather
union
select mother, child
from isMother
```

```
select father, child
from isFather
union
select child, mother
from isMother
```

These queries are different!



# Positional notation

With renaming (same as before):

```
select father as parent, child
from   isFather
union
select child, mother as parent
from   isMother
```

If we want to return fathers and mothers as parents, this is the correct query:

```
select father as parent, child
from   isFather
union
select mother as parent, child
from   isMother
```



## Difference

```
select name
from   employee
except
select surname as name
from   employee
```

Note: **except** eliminates duplicate tuples

Note: **except all** does not eliminate duplicate tuples

The difference can also be expressed by nested **select** statements.



# Intersection

```
select name
from   employee
intersect
select surname as name
from   employee
```

is equivalent to

```
select distinct i.name
from   employee i, employee j
where  i.name = j.surname
```

Note: **intersect** eliminates duplicate tuples

Note: **intersect all** does not eliminate duplicate tuples



## Nested queries

- A nested **select** statement can appear as a condition in the where clause
- In particular, the conditions allow for:
  - comparing an attribute (or a sequence of attributes) with the result of a sub-query
  - existential quantification



## Nested queries: example

name and income of Franco's father:

```
select  name, income
from    person, isFather
where   name = father and child = 'Franco'
```

```
select  name, income
from    person
where   name = (select father
                  from    isFather
                  where   child = 'Franco')
```



## Nested queries: operators

The result of a nested query can be compared in the **where** clause using several **operators**:

- Equality and the other comparison operator: in this case, the result of the nested query must be a single tuple
- If the result of the nested query may contain multiple tuples, the nested query can be preceded by:
  - **any**: returns true if the comparison is true for **at least** one of the tuples in the result of the nested query
  - **all**: returns true if the comparison is true for **every** tuple in the result of the nested query
- The operator **in**, which is equivalent to **=any**
- The operator **not in**, which is equivalent to **<>all**
- The operator **exists**



## Nested queries: example

name and income of the fathers of persons earning more than 20 millions:

```
select distinct p.name, p.income
from   person p, isFather, person c
where  p.name = father and child = c.name
       and c.income > 20
```

```
select name, income
from   person
where  name = any
```

```
(select father
   from   isFather, person
  where  child = name
        and income > 20)
```

Fathers of persons  
earning more than  
20 millions



## Nested queries: example

name and income of the fathers of persons earning more than 20 millions:

```
select name, income
from person
where name in (select father
                from isFather, person
                where child = name
                and income > 20)
```

```
select name, income
from person
where name in (select father
                from isFather
                where child in (select name
                                from person
                                where income > 20)
```

Fathers of persons earning more than 20 millions

Persons earning more than 20 millions

(select name  
from person  
where income > 20)



## Nested queries: all (example)

Persons whose income is greater than the income of every person who is less than 30 years old:

```
select name
from   person
where  income > all ( select income
                      from person
                      where age < 30 )
```



## Nested queries: exists (example)

The **exists** operator is used to return true if the results of the sub-query is **not empty**.

*Example:* persons having at least a child.

```
select *  
from   person p  
where  exists (select *  
                from   isFather  
                where  father = p.name)  
       or  
       exists (select *  
                from   isMother  
                where  mother = p.name)
```

Notice that the attribute **name** refers to the relation in the **from** clause.



## Exercise 9: nested queries

Return name and age of mothers having at least a child who is less than 18 years old.



## Exercise 9: nested queries

Return name and age of mothers having at least a child who is less than 18 years old.

**Solution 1:** a join to select name and age of mothers, and a sub-query for the condition on the children

**Solution 2:** two sub-queries and no join



## Exercise 9: solution 1

Return name and age of mothers having at least a child who is less than 18 years old.

```
select name, age
from   person, isMother
where  name = mother and
       child in (select name
                  from   person
                  where  age < 18)
```



## Exercise 9: solution 2

Return name and age of mothers having at least a child who is less than 18 years old.

```
select name, age
from person
where name in (select mother
                from isMother
                where child in (select name
                                from person
                                where age<18))
```





## Nested queries: comments

- Nested queries may pose performance problems to the DBMSs (since they are not very good in optimizing the execution of such statements)
- However, nested queries are sometimes more readable than equivalent, non-nested ones.
- In some systems, sub-queries cannot contain set operators, but this is not a significant limitation.



## Nested queries, comments

- **visibility** rules:
  - It is not possible to refer to variables (attributes) defined in inner blocks
  - If a variable or table name is omitted, the assumption is that it refers to the «closest» variable or table
- A block can refer to variables defined in the same block or in outer blocks, unless they are hidden by definitions of variables with the same name.
- **Semantics**: the inner query is executed once **for each tuple** of the outer query



## Nested queries: visibility

Return the persons having at least a child.

```
select *  
from   person  
where  exists (select *  
                from   isFather  
                where  father = name)  
or  
exists (select *  
        from   isMother  
        where  mother = name)
```

Attribute **name** refers to the relation **person** in the **from** clause.



## Nested queries: visibility

The following query is **incorrect**:

```
select *  
from employee  
where depart in (select name  
                  from department D1  
                  where name = 'Produzione')  
or  
   depart in (select name  
               from department D2  
               where D2.citta = D1.citta)
```

**employee**

name	surname	depart
------	---------	--------

**department**

name	address	city
------	---------	------

## Example

name and income of the fathers of persons earning more than 20 millions,  
**returning the child's income too.**

```
select distinct p.name, p.income, c.income
from   person p, isFather, person c
where  p.name = father and child = c.name
       and c.income > 20
```

In this case the following “intuitive” nested query **is not correct**:

```
select name, income, c.income
from person
where name in (select father
               from isFather
               where child in (select name
                              from person c
                              where c.income > 20))
```



## Nested and correlated queries

It may be necessary to use, in an inner block, variables defined in outer blocks: in this case the query is called nested and **correlated**.

*Example:* fathers whose children earn more than 20 millions.

```
select distinct father
from isFather z
where not exists (select *
                  from isFather w, person
                  where w.father = z.father
                        and w.child = name
                        and income <= 20)
```



## Exercise 10: nested and correlated queries

Return name and age of every mother having at least a child who is less than 30 years younger than her.



## Exercise 10: solution

Return name and age of every mother having at least a child who is less than 30 years younger than her.

```
select name, age
from   person p, isMother
where  name = mother and
       child in (select name
                  from   person
                  where  p.age - age < 30)
```





## Difference can be expressed by nested queries

```
select name from employee
```

```
except
```

```
select surname as name from employee
```

```
select name
```

```
from employee
```

```
where name not in (select surname  
                    from employee)
```



## Intersection can be expressed by nested queries

```
select name from employee
```

**intersection**

```
select surname from employee
```

```
select name
```

```
from employee
```

```
where name in (select surname  
               from employee)
```



## Exercise 11: nesting and functions

Return the person(s) with maximum income.



## Exercise 11: solution

Return the person(s) with maximum income.

```
select *  
from person  
where income = (select max(income)  
                from person)
```

or:

```
select *  
from person  
where income >= all (select income  
                    from person)
```



## Nested queries: condition on multiple attributes

Return the persons whose pair (age, income) is different from all other persons.

```
select *  
from person p  
where (age,income) not in  
      (select age, income  
       from person  
       where name <> p.name)
```



## Nested queries in the **from** clause

Nested queries may appear not only in the **where** clause, but also in the **from** clause:

```
select p.father
from isFather p, (select name
                  from person
                  where age > 30) c
where c.name = p.child
```

Semantics: the table whose alias is **f**, and defined as a nested query in the **from** clause, is not a database table, but is computed using the associated **select** query.



# SQL

## 4. Further aspects

1. Data definition
2. Data manipulation
3. Queries
- 4. Further aspects**



# Generic integrity constraints: check

To specify complex constraints on a tuple or a table:

`check` (*Condition*)

```
create table employee
( ID character(6) ,
  surname character(20) ,
  name character(20) ,
  sex character not null check (sex in ('M' , 'F'))
  salary integer,
  manager character(6) ,
  check (salary <= (select salary
                      from   employee j
                      where  manager = j.ID))
)
```





# Views

- A view is a table **whose instance is derived from other tables through a query.**

```
create view ViewName [(AttributeList)] as SelectSQL
```

- Views are virtual tables: their instance is computed only when they are used by other queries.
- *Example:*

```
create view adminEmp (Mat, name, surname, Stip) as  
  select ID, name, surname, salary  
  from   employee  
  where  Depart = 'Administration' and  
         salary > 10
```

## Nested queries in the **having** clause

- Return the age of persons such that sum of the income of persons having that age is maximum.
- Assuming there are no null values in the income attribute, and using a nested query in the **having** clause:

```
select age
from   person
group by age
having sum(income) >= all (select sum(income)
                           from person
                           group by age)
```



## Solution with views

```
create view ageincome(age,total-income) as  
  select age, sum(income)  
  from    person  
  group by age
```

```
select age  
from    ageincome  
where   total-income = (select max(total-income)  
                        from ageincome)
```